

```
In [1]: import pandas as pd
from statsbombpy import sb
from mplsoccer import Pitch
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import networkx as nx
from tqdm import tqdm
from sklearn.preprocessing import LabelEncoder
```

```
In [2]: df=sb.competitions()
```

```
In [3]: df = df.sort_values('competition_id')
```

```
In [4]: df.head(60)
```

Out[4]:

	competition_id	season_id	country_name	competition_name	competition_gender
64	2	27	England	Premier League	male
65	2	44	England	Premier League	male
58	7	235	France	Ligue 1	male
59	7	108	France	Ligue 1	male
60	7	27	France	Ligue 1	male
0	9	281	Germany	1. Bundesliga	male
1	9	27	Germany	1. Bundesliga	male
48	11	22	Spain	La Liga	male
55	11	278	Spain	La Liga	male
47	11	23	Spain	La Liga	male
54	11	37	Spain	La Liga	male
53	11	38	Spain	La Liga	male
52	11	39	Spain	La Liga	male
51	11	40	Spain	La Liga	male
50	11	41	Spain	La Liga	male
49	11	21	Spain	La Liga	male
39	11	42	Spain	La Liga	male
40	11	4	Spain	La Liga	male
41	11	1	Spain	La Liga	male
42	11	2	Spain	La Liga	male
43	11	27	Spain	La Liga	male

competition_id	season_id	country_name	competition_name	competition_gender	competition_type
44	11	26	Spain	La Liga	male
45	11	25	Spain	La Liga	male
46	11	24	Spain	La Liga	male
38	11	90	Spain	La Liga	male
67	12	86	Italy	Serie A	male
66	12	27	Italy	Serie A	male
3	16	4	Europe	Champions League	male
6	16	27	Europe	Champions League	male
15	16	37	Europe	Champions League	male
7	16	26	Europe	Champions League	male
8	16	25	Europe	Champions League	male
4	16	1	Europe	Champions League	male
20	16	276	Europe	Champions League	male
19	16	71	Europe	Champions League	male
18	16	277	Europe	Champions League	male
17	16	76	Europe	Champions League	male
16	16	44	Europe	Champions League	male
5	16	2	Europe	Champions League	male
14	16	39	Europe	Champions League	male
13	16	41	Europe	Champions League	male
12	16	21	Europe	Champions League	male

competition_id	season_id	country_name	competition_name	competition_gender	...
11	16	22	Europe	Champions League	male
10	16	23	Europe	Champions League	male
9	16	24	Europe	Champions League	male
70	35	75	Europe	UEFA Europa League	male
26	37	42	England	FA Women's Super League	female
25	37	90	England	FA Women's Super League	female
27	37	4	England	FA Women's Super League	female
36	43	269	International	FIFA World Cup	male
35	43	270	International	FIFA World Cup	male
34	43	272	International	FIFA World Cup	male
33	43	51	International	FIFA World Cup	male
32	43	54	International	FIFA World Cup	male
31	43	55	International	FIFA World Cup	male
30	43	3	International	FIFA World Cup	male
29	43	106	International	FIFA World Cup	male
61	44	107	United States of America	Major League Soccer	male
63	49	3	United States of America	NWSL	female
71	53	106	Europe	UEFA Women's Euro	female

In [5]:

```
# Specify competition_id
competition_id = 11 # Replace with your desired competition ID

# Fetch all available seasons for the competition
competitions = sb.competitions()
seasons = competitions[competitions['competition_id'] == competition_id]['season_id']
```

```
# List to store all matches and events
all_events = []

# Loop through each season and fetch matches and events
for season_id in tqdm(seasons, desc="Processing Seasons"):
    try:
        # Fetch matches for the current season
        matches = sb.matches(competition_id=competition_id, season_id=season_id)

        # Loop through matches and fetch events
        for match_id in matches['match_id']:
            try:
                # Fetch events for the current match
                events = sb.events(match_id=match_id)

                # Add match_id, home_team, and away_team to the events DataFrame
                home_team = matches[matches['match_id'] == match_id]['home_team']
                away_team = matches[matches['match_id'] == match_id]['away_team']

                events['match_id'] = match_id
                events['home_team'] = home_team
                events['away_team'] = away_team

                # Append the events to the list
                all_events.append(events)
            except Exception as e:
                print(f"Error fetching events for match_id {match_id}: {e}")
    except Exception as e:
        print(f"Error fetching matches for season_id {season_id}: {e}")

# Combine all events into a single DataFrame
if all_events:
    e_df = pd.concat(all_events, ignore_index=True)
    print(e_df.head())
else:
    print("No events data fetched.")
```

Processing Seasons: 100%|██████████| 18/18 [09:03<00:00, 30.19s/it]

```

50_50 bad_behaviour_card ball_receipt_outcome ball_recovery_offensive \
0   NaN           NaN           NaN           NaN
1   NaN           NaN           NaN           NaN
2   NaN           NaN           NaN           NaN
3   NaN           NaN           NaN           NaN
4   NaN           NaN           NaN           NaN

ball_recovery_failure carry_end_location clearance_aerial_won \
0                   NaN           NaN           NaN
1                   NaN           NaN           NaN
2                   NaN           NaN           NaN
3                   NaN           NaN           NaN
4                   NaN           NaN           NaN

clearance_body_part clearance_head clearance_left_foot ... shot_redirect \
0           NaN           NaN           NaN   ...
1           NaN           NaN           NaN   ...
2           NaN           NaN           NaN   ...
3           NaN           NaN           NaN   ...
4           NaN           NaN           NaN   ...

shot_follows_dribble goalkeeper_success_in_play goalkeeper_lost_in_play \
0           NaN           NaN           NaN
1           NaN           NaN           NaN
2           NaN           NaN           NaN
3           NaN           NaN           NaN
4           NaN           NaN           NaN

half_start_late_video_start player_off_permanent goalkeeper_lost_out \
0           NaN           NaN           NaN
1           NaN           NaN           NaN
2           NaN           NaN           NaN
3           NaN           NaN           NaN
4           NaN           NaN           NaN

goalkeeper_success_out half_end_early_video_end goalkeeper_saved_to_post
0           NaN           NaN           NaN
1           NaN           NaN           NaN
2           NaN           NaN           NaN
3           NaN           NaN           NaN
4           NaN           NaN           NaN

```


[5 rows x 121 columns]

In [6]: e_df

Out[6]:

	50_50	bad_behaviour_card	ball_receipt_outcome	ball_recovery_offensive	ball
0	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN
...
3136478	NaN	NaN	NaN	NaN	NaN
3136479	NaN	NaN	NaN	NaN	NaN
3136480	NaN	NaN	NaN	NaN	NaN
3136481	NaN	NaN	NaN	NaN	NaN
3136482	NaN	NaN	NaN	NaN	NaN

3136483 rows × 121 columns



In [7]:

```
# Filter shot events
shots = e_df[e_df['type'] == 'Shot']
```

In [8]:

```
# Initialize home and away scores and game_state in the shots DataFrame using .loc
shots.loc[:, 'home_score'] = 0
shots.loc[:, 'away_score'] = 0
shots.loc[:, 'game_state'] = 'tied' # Default game state is 'tied'

# Loop through each match and update the game_state for each shot event in the s
for match_id in shots['match_id'].unique():
    # Extract shot events for the current match
    match_shots = shots[shots['match_id'] == match_id]

    home_goals = 0
    away_goals = 0

    # Loop through the shot events for this match to calculate game state
    for i, event in match_shots.iterrows():
        # Update scores if a goal is scored
        if event['shot_outcome'] == 'Goal':
            if event['team'] == event['home_team']:
                home_goals += 1
            elif event['team'] == event['away_team']:
                away_goals += 1

        # Set the current home and away scores for the shot event using .loc
        shots.loc[i, 'home_score'] = home_goals
        shots.loc[i, 'away_score'] = away_goals

        # Determine the game state based on the team taking the shot
        if event['team'] == event['home_team']:
            if home_goals > away_goals:
                shots.loc[i, 'game_state'] = 'leading'
```

```

        elif home_goals < away_goals:
            shots.loc[i, 'game_state'] = 'trailing'
        else:
            shots.loc[i, 'game_state'] = 'tied'

    elif event['team'] == event['away_team']:
        if away_goals > home_goals:
            shots.loc[i, 'game_state'] = 'leading'
    elif away_goals < home_goals:
        shots.loc[i, 'game_state'] = 'trailing'
    else:
        shots.loc[i, 'game_state'] = 'tied'

# Display the updated shots DataFrame with game_state
print(shots[['match_id', 'home_team', 'away_team', 'home_score', 'away_score', 'game_state']])

```

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2881702795.py:2: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe = frame.copy()`

shots.loc[:, 'home_score'] = 0

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2881702795.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

shots.loc[:, 'home_score'] = 0

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2881702795.py:3: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe = frame.copy()`

shots.loc[:, 'away_score'] = 0

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2881702795.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

shots.loc[:, 'away_score'] = 0

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2881702795.py:4: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe = frame.copy()`

shots.loc[:, 'game_state'] = 'tied' # Default game state is 'tied'

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2881702795.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

shots.loc[:, 'game_state'] = 'tied' # Default game state is 'tied'

	match_id	home_team	away_team	home_score	away_score	game_state
3805	3773386	Deportivo Alavés	Barcelona	0	0	tied
3806	3773386	Deportivo Alavés	Barcelona	0	0	tied
3807	3773386	Deportivo Alavés	Barcelona	0	0	tied
3808	3773386	Deportivo Alavés	Barcelona	0	0	tied
3809	3773386	Deportivo Alavés	Barcelona	1	0	leading

```
In [9]: # Ensure 'timestamp' is in datetime format if it's not already
shots.loc[:, 'timestamp'] = pd.to_datetime(shots['timestamp'])

# Initialize the 'time_since_last_event' column to NaN using .loc
shots.loc[:, 'time_since_last_event'] = None

# Loop through each match to calculate the time difference
for match_id in shots['match_id'].unique():
    # Extract the shots for the current match
    match_shots = shots[shots['match_id'] == match_id]

    # Loop through the shots to calculate the time since the last event
    for i in range(1, len(match_shots)):
        current_event_time = match_shots.iloc[i]['timestamp']
        previous_event_time = match_shots.iloc[i - 1]['timestamp']

        # Calculate the time difference in seconds (you can change this to minutes)
        time_diff = (current_event_time - previous_event_time).total_seconds()

        # Update the 'time_since_last_event' for the current event using .loc
        shots.loc[match_shots.index[i], 'time_since_last_event'] = time_diff

# Display the updated shots DataFrame with 'time_since_last_event'
print(shots[['match_id', 'timestamp', 'time_since_last_event']].head())
```

C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2026288433.py:5: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe = frame.copy()`
shots.loc[:, 'time_since_last_event'] = None
C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\2026288433.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
shots.loc[:, 'time_since_last_event'] = None
   match_id          timestamp time_since_last_event
3805    3773386 2025-02-11 00:12:19.834000            None
3806    3773386 2025-02-11 00:16:07.385000         227.551
3807    3773386 2025-02-11 00:19:42.833000         215.448
3808    3773386 2025-02-11 00:22:26.180000         163.347
3809    3773386 2025-02-11 00:30:06.585000        460.405
```

```
In [10]: # Function to calculate 2D or 3D distance based on the shot_end_location type
def calculate_shot_distance(row):
    if len(row['shot_end_location']) == 3:
        return np.sqrt(
            (row['location'][0] - row['shot_end_location'][0]) ** 2 +
            (row['location'][1] - row['shot_end_location'][1]) ** 2 +
            (0 - row['shot_end_location'][2]) ** 2
```

```

        )
else:
    return np.sqrt(
        (row['location'][0] - row['shot_end_location'][0]) ** 2 +
        (row['location'][1] - row['shot_end_location'][1]) ** 2
    )

# Function to categorize shot placement into sections of the goal (Top-Corner, B
def categorize_goal_section(placement, y_value):
    # Left or right top/bottom corners (y-value is within 36-37 or 43-44)
    if 36 <= y_value <= 37 or 43 <= y_value <= 44:
        if placement <= 0.5: # Bottom-Corner
            return 'Bottom-Corner'
        elif placement > 2.2: # Top-Corner
            return 'Top-Corner'

    # Center section of the goal (y-value is between 37 and 43)
    elif 37 < y_value < 43:
        if placement <= 0.5: # Bottom
            return 'Bottom'
        elif placement > 2.2: # Top
            return 'Top'
        else: # Middle
            return 'Middle'

    # Ensure that shots falling in the middle section (placement between 1.0 and
    elif 1.0 < placement <= 2.5: # Middle section
        return 'Middle'

    # For shots outside the valid placement, categorize as off-target
    return 'Off-Target'

# Function to check if a shot is off-target based on x, y, and z ranges
def is_off_target(row):
    valid_x_range = (0, 120)
    valid_y_range = (36, 44)
    valid_z_range = (0, 2.67)

    # Check if x value is out of valid range
    if not (valid_x_range[0] <= row['shot_end_location'][0] <= valid_x_range[1]):
        return True

    # Check if y value is out of valid range
    if not (valid_y_range[0] <= row['shot_end_location'][1] <= valid_y_range[1]):
        return True

    # Check if z value is out of valid range, unless z = 0
    if len(row['shot_end_location']) == 3 and not (valid_z_range[0] <= row['shot_end_location'][2] <= valid_z_range[1]):
        # Only treat as off-target if z is out of range (except for z=0, which is
        if row['shot_end_location'][2] != 0:
            return True

    # If all checks pass, the shot is on target
    return False

# Function to process the shot and add additional features
def process_shot(shots):
    shots = shots.copy()

    # Calculate shot distance based on start and end location

```

```

shots['shot_distance'] = shots.apply(calculate_shot_distance, axis=1)

# Calculate shot angle in radians and convert to degrees
shots['shot_angle'] = np.arctan2(
    shots['shot_end_location'].apply(lambda loc: loc[1] if len(loc) == 2 else 0),
    shots['shot_end_location'].apply(lambda loc: loc[0] if len(loc) == 2 else 1))
shots['shot_angle_degrees'] = np.degrees(shots['shot_angle'])

# Handle Z-value explicitly when shot_end_location has 2 dimensions (set Z to 0)
shots['shot_placement'] = shots['shot_end_location'].apply(lambda loc: loc[2] if len(loc) == 3 else 0)

# Categorize shot placement into goal sections (Top-Corner, Bottom-Corner, Top-Middle, etc)
shots['goal_section'] = shots.apply(lambda row: categorize_goal_section(row))

# Apply the off-target check
shots['off_target'] = shots.apply(is_off_target, axis=1)

# Final classification: If the shot is off-target, label as 'Off-Target', otherwise as 'On-Target'
shots['shot_classification'] = np.where(shots['off_target'], 'Off-Target', 'On-Target')

return shots

shots = process_shot(shots)

```

In [11]: `print(shots.dtypes)`

50_50	object
bad_behaviour_card	object
ball_receipt_outcome	object
ball_recovery_offensive	object
ball_recovery_recovery_failure	object
...	
shot_angle_degrees	float64
shot_placement	float64
goal_section	object
off_target	bool
shot_classification	object

Length: 132, dtype: object

In [12]: `# List of important columns you want to include in your model`

```

important_columns = [
    'match_id', 'timestamp', 'location', 'shot_end_location', 'shot_first_time',
    'shot_technique', 'shot_body_part', 'shot_type', 'shot_outcome', 'under_pressure',
    'shot_distance', 'shot_angle_degrees', 'time_since_last_event', 'game_state',
    'home_score', 'away_score', 'position', 'period', 'shot_classification']

```

```

refined_model_data = shots[important_columns].copy()

refined_model_data.head(50)

```

Out[12]:		match_id	timestamp	location	shot_end_location	shot_first_time	shot_techn
3805	3773386	2025-02-11 00:12:19.834000	[108.6, 28.0]	[120.0, 47.8, 0.0]	NaN	Nc	
3806	3773386	2025-02-11 00:16:07.385000	[103.6, 51.0]	[115.8, 42.1, 0.0]	True	Nc	
3807	3773386	2025-02-11 00:19:42.833000	[104.3, 33.9]	[120.0, 50.6, 0.0]	True	Nc	
3808	3773386	2025-02-11 00:22:26.180000	[97.9, 44.3]	[119.8, 37.6]	NaN	Nc	
3809	3773386	2025-02-11 00:30:06.585000	[118.3, 42.1]	[120.0, 40.7, 0.0]	NaN	Nc	
3810	3773386	2025-02-11 00:35:25.546000	[113.1, 39.3]	[120.0, 38.5, 4.9]	NaN	Nc	
3811	3773386	2025-02-11 00:42:34.293000	[101.3, 41.6]	[110.5, 39.5]	NaN	Nc	
3812	3773386	2025-02-11 00:44:47.440000	[84.7, 42.1]	[120.0, 46.4, 5.6]	NaN	Nc	
3813	3773386	2025-02-11 00:45:16.926000	[105.3, 32.9]	[118.3, 37.0, 0.0]	True	Nc	
3814	3773386	2025-02-11 00:02:47.901000	[113.3, 26.1]	[118.2, 35.8, 1.3]	NaN	Nc	
3815	3773386	2025-02-11 00:05:54.058000	[107.4, 27.2]	[117.1, 36.3, 0.5]	NaN	Nc	
3816	3773386	2025-02-11 00:08:34.260000	[111.5, 31.6]	[120.0, 47.7, 0.7]	True	V	
3817	3773386	2025-02-11 00:14:25.864000	[99.0, 24.7]	[120.0, 34.0, 5.3]	NaN	Nc	
3818	3773386	2025-02-11 00:17:57.784000	[107.6, 31.3]	[120.0, 42.7, 0.4]	True		
3819	3773386	2025-02-11 00:19:34.950000	[108.4, 49.5]	[115.2, 43.3, 0.3]	NaN	Nc	
3820	3773386	2025-02-11 00:22:39.838000	[98.2, 34.7]	[120.0, 25.8, 0.2]	True	Nc	
3821	3773386	2025-02-11 00:23:34.135000	[106.6, 35.3]	[117.7, 39.8, 1.6]	NaN	Nc	

	match_id	timestamp	location	shot_end_location	shot_first_time	shot_techn
3822	3773386	2025-02-11 00:28:44.317000	[103.2, 35.7]	[104.1, 36.0]	NaN	Half V
3823	3773386	2025-02-11 00:32:00.879000	[103.4, 24.7]	[116.5, 37.6, 2.0]	NaN	Nc
3824	3773386	2025-02-11 00:32:50.235000	[96.4, 46.8]	[98.2, 46.0]	NaN	Nc
3825	3773386	2025-02-11 00:37:44.138000	[97.5, 49.6]	[99.7, 48.4]	NaN	Nc
3826	3773386	2025-02-11 00:38:05.308000	[92.9, 28.0]	[102.5, 34.3]	NaN	Nc
3827	3773386	2025-02-11 00:42:02.502000	[106.4, 49.8]	[117.8, 41.2, 0.9]	True	Nc
3828	3773386	2025-02-11 00:42:38.083000	[88.7, 39.7]	[116.6, 39.3, 2.1]	NaN	Nc
3829	3773386	2025-02-11 00:43:25.053000	[114.5, 46.0]	[118.2, 40.8]	NaN	Nc
3830	3773386	2025-02-11 00:44:08.947000	[108.9, 44.1]	[120.0, 27.2, 0.0]	True	Nc
3831	3773386	2025-02-11 00:45:44.395000	[94.0, 49.1]	[100.2, 47.5]	NaN	Nc
3832	3773386	2025-02-11 00:48:11.821000	[99.0, 46.8]	[107.8, 45.0]	True	Nc
3833	3773386	2025-02-11 00:48:33.533000	[104.9, 49.8]	[120.0, 33.4, 1.9]	NaN	Nc
7516	3773565	2025-02-11 00:01:39.034000	[101.1, 50.9]	[117.2, 41.6, 1.6]	NaN	Nc
7517	3773565	2025-02-11 00:10:28.938000	[106.6, 34.6]	[107.6, 34.9]	NaN	Nc
7518	3773565	2025-02-11 00:11:06.691000	[110.4, 37.9]	[120.0, 43.4, 0.2]	NaN	Nc
7519	3773565	2025-02-11 00:15:41.304000	[103.1, 57.1]	[120.0, 33.8, 0.2]	NaN	Nc
7520	3773565	2025-02-11 00:16:47.548000	[92.6, 47.5]	[94.7, 47.1]	NaN	Nc
7521	3773565	2025-02-11 00:22:23.374000	[109.5, 55.4]	[120.0, 25.1, 0.1]	True	Nc

	match_id	timestamp	location	shot_end_location	shot_first_time	shot_techn
7522	3773565	2025-02-11 00:29:17.718000	[103.1, 31.4]	[104.2, 31.8]	NaN	Nc
7523	3773565	2025-02-11 00:32:38.713000	[99.2, 56.5]	[117.9, 43.2, 0.4]	NaN	Nc
7524	3773565	2025-02-11 00:34:58.025000	[103.9, 37.9]	[120.0, 36.6, 2.1]	NaN	Nc
7525	3773565	2025-02-11 00:38:23.131000	[100.5, 45.8]	[120.0, 40.9, 7.6]	True	Nc
7526	3773565	2025-02-11 00:41:54.600000	[100.8, 43.8]	[120.0, 36.1, 0.2]	NaN	Nc
7527	3773565	2025-02-11 00:04:55.141000	[104.0, 35.9]	[117.4, 38.4, 0.2]	NaN	Nc
7528	3773565	2025-02-11 00:10:47.410000	[102.1, 45.6]	[120.0, 37.7, 4.0]	NaN	Nc
7529	3773565	2025-02-11 00:14:25.059000	[111.1, 32.3]	[120.0, 38.6, 5.5]	True	Half V
7530	3773565	2025-02-11 00:17:58.795000	[113.2, 52.5]	[120.0, 36.5, 0.2]	NaN	Half V
7531	3773565	2025-02-11 00:21:18.935000	[107.6, 29.1]	[120.0, 44.3, 3.4]	NaN	Half V
7532	3773565	2025-02-11 00:29:15.684000	[103.5, 49.5]	[120.0, 43.4, 7.6]	True	Half V
7533	3773565	2025-02-11 00:34:14.506000	[100.2, 33.3]	[107.8, 35.0]	NaN	Nc
7534	3773565	2025-02-11 00:34:50.223000	[94.0, 44.4]	[101.8, 43.9]	NaN	Nc
7535	3773565	2025-02-11 00:35:57.390000	[100.6, 58.9]	[116.1, 42.1, 0.7]	NaN	Nc
7536	3773565	2025-02-11 00:39:45.616000	[103.8, 56.4]	[106.8, 53.9]	NaN	Nc

In [13]: `# Check for missing values in the important columns
missing_values = refined_model_data[important_columns].isnull().sum()`

In [14]: `missing_values`

```
Out[14]: match_id          0
         timestamp        0
         location         0
         shot_end_location 0
         shot_first_time    14505
         shot_technique      0
         shot_body_part       0
         shot_type          0
         shot_outcome         0
         under_pressure      17448
         shot_distance        0
         shot_angle_degrees   0
         time_since_last_event 868
         game_state          0
         home_score          0
         away_score          0
         position            0
         period              0
         shot_classification   0
         dtype: int64
```

```
In [15]: cleaned_model_data = refined_model_data.copy()
```

```
# Assuming `df` is your original dataframe
cleaned_model_data['time_since_last_event'] = pd.to_numeric(refined_model_data['time_since_last_event'])

# Fill NaN values with the mean (you can also use median or any other method you prefer)
cleaned_model_data['time_since_last_event'].fillna(refined_model_data['time_since_last_event'].mean(), inplace=True)

# Display the cleaned data
cleaned_model_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 21210 entries, 3805 to 3136364
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   match_id         21210 non-null   int64  
 1   timestamp        21210 non-null   object  
 2   location         21210 non-null   object  
 3   shot_end_location 21210 non-null   object  
 4   shot_first_time    6705 non-null   object  
 5   shot_technique      21210 non-null   object  
 6   shot_body_part       21210 non-null   object  
 7   shot_type          21210 non-null   object  
 8   shot_outcome        21210 non-null   object  
 9   under_pressure      3762 non-null   object  
 10  shot_distance       21210 non-null   float64 
 11  shot_angle_degrees 21210 non-null   float64 
 12  time_since_last_event 21210 non-null   float64 
 13  game_state         21210 non-null   object  
 14  home_score          21210 non-null   int64  
 15  away_score          21210 non-null   int64  
 16  position            21210 non-null   object  
 17  period              21210 non-null   int64  
 18  shot_classification 21210 non-null   object  
dtypes: float64(3), int64(4), object(12)
memory usage: 3.7+ MB
```

```
C:\Users\harsh\AppData\Local\Temp\ipykernel_12040\3237411695.py:8: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
cleaned_model_data['time_since_last_event'].fillna(refined_model_data['time_since_last_event'].mean(), inplace=True)
```

```
In [16]: # Function to convert NaN and non-boolean values to boolean
def to_boolean(x):
    if pd.isna(x): # If the value is NaN
        return False
    return bool(x)

# Apply the function to convert to boolean
cleaned_model_data['shot_first_time'] = cleaned_model_data['shot_first_time'].apply(to_boolean)
cleaned_model_data['under_pressure'] = cleaned_model_data['under_pressure'].apply(to_boolean)
```

```
In [17]: cleaned_model_data.info()
```

#	Column	Non-Null Count	Dtype
0	match_id	21210	non-null int64
1	timestamp	21210	non-null object
2	location	21210	non-null object
3	shot_end_location	21210	non-null object
4	shot_first_time	21210	non-null bool
5	shot_technique	21210	non-null object
6	shot_body_part	21210	non-null object
7	shot_type	21210	non-null object
8	shot_outcome	21210	non-null object
9	under_pressure	21210	non-null bool
10	shot_distance	21210	non-null float64
11	shot_angle_degrees	21210	non-null float64
12	time_since_last_event	21210	non-null float64
13	game_state	21210	non-null object
14	home_score	21210	non-null int64
15	away_score	21210	non-null int64
16	position	21210	non-null object
17	period	21210	non-null int64
18	shot_classification	21210	non-null object

dtypes: bool(2), float64(3), int64(4), object(10)
memory usage: 3.5+ MB

```
In [18]: from sklearn.preprocessing import StandardScaler

# Initialize StandardScaler
scaler = StandardScaler()

# List of numerical columns to scale
```

```

numerical_columns = ['shot_distance', 'shot_angle_degrees']

# Apply StandardScaler to the numerical columns
cleaned_model_data[numerical_columns] = scaler.fit_transform(cleaned_model_data[

# Check the scaled data
cleaned_model_data[numerical_columns]

```

Out[18]:

	shot_distance	shot_angle_degrees
3805	0.716722	0.821318
3806	-0.003108	0.290491
3807	0.723585	1.170489
3808	0.721800	-0.488807
3809	-1.201807	-0.089012
...
3136360	-1.006759	1.544207
3136361	-0.344708	0.170933
3136362	0.605622	0.087028
3136363	-0.859126	-0.770692
3136364	-1.079093	1.797731

21210 rows × 2 columns

In [19]:

```
# Drop the columns that are not useful for the model
cleaned_model_data = cleaned_model_data.drop(columns=['timestamp', 'match_id', 'l
```

In [20]:

```
cleaned_model_data['Goal'] = cleaned_model_data['shot_outcome'].apply(lambda x:
cleaned_model_data.drop(columns=['shot_outcome'], inplace=True)
```

In [21]:

```
# Map integers to meaningful Labels
period_mapping = {
    1: "First Half",
    2: "Second Half",
    3: "Third Period",
    4: "Fourth Period",
    5: "Penalties"
}

# Apply the mapping to the 'period' column
cleaned_model_data['period'] = cleaned_model_data['period'].map(period_mapping)

# Convert the column to a categorical type
cleaned_model_data['period'] = pd.Categorical(
    cleaned_model_data['period'],
    categories=["First Half", "Second Half", "Third Period", "Fourth Period", "P
    ordered=True
)
```

```
In [22]: catboost_data=cleaned_model_data.copy()
```

```
In [23]: # Apply One-Hot Encoding using pandas get_dummies
categorical_columns = ['shot_technique', 'shot_body_part', 'shot_type', 'game_st


# Perform One-Hot Encoding on the specified columns
cleaned_model_data = pd.get_dummies(cleaned_model_data, columns=categorical_colu

# Check the changes
cleaned_model_data
```

Out[23]:

	shot_first_time	under_pressure	shot_distance	shot_angle_degrees	time_since...
3805	False	False	0.716722	0.821318	1
3806	True	False	-0.003108	0.290491	2
3807	True	False	0.723585	1.170489	2
3808	False	False	0.721800	-0.488807	1
3809	False	False	-1.201807	-0.089012	2
...
3136360	False	False	-1.006759	1.544207	2
3136361	False	True	-0.344708	0.170933	2
3136362	False	False	0.605622	0.087028	2
3136363	True	False	-0.859126	-0.770692	2
3136364	False	False	-1.079093	1.797731	

21210 rows × 62 columns



```
In [24]: cleaned_model_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 21210 entries, 3805 to 3136364
Data columns (total 62 columns):
 #   Column           Non-Null Count Dtype
 ---  -----
 0   shot_first_time    21210 non-null  bool
 1   under_pressure     21210 non-null  bool
 2   shot_distance      21210 non-null  float64
 3   shot_angle_degrees 21210 non-null  float64
 4   time_since_last_event 21210 non-null  float64
 5   home_score          21210 non-null  int64
 6   away_score          21210 non-null  int64
 7   Goal                21210 non-null  int64
 8   shot_technique_Backheel 21210 non-null  bool
 9   shot_technique_Diving Header 21210 non-null  bool
 10  shot_technique_Half Volley 21210 non-null  bool
 11  shot_technique_Lob       21210 non-null  bool
 12  shot_technique_Normal 21210 non-null  bool
 13  shot_technique_Overhead Kick 21210 non-null  bool
 14  shot_technique_Volley   21210 non-null  bool
 15  shot_body_part_Head    21210 non-null  bool
 16  shot_body_part_Left Foot 21210 non-null  bool
 17  shot_body_part_Other   21210 non-null  bool
 18  shot_body_part_Right Foot 21210 non-null  bool
 19  shot_type_Corner      21210 non-null  bool
 20  shot_type_Free Kick   21210 non-null  bool
 21  shot_type_Open Play   21210 non-null  bool
 22  shot_type_Penalty     21210 non-null  bool
 23  game_state_leading   21210 non-null  bool
 24  game_state_tied      21210 non-null  bool
 25  game_state_trailing  21210 non-null  bool
 26  position_Center Attacking Midfield 21210 non-null  bool
 27  position_Center Back   21210 non-null  bool
 28  position_Center Defensive Midfield 21210 non-null  bool
 29  position_Center Forward 21210 non-null  bool
 30  position_Center Midfield 21210 non-null  bool
 31  position_Goalkeeper   21210 non-null  bool
 32  position_Left Attacking Midfield 21210 non-null  bool
 33  position_Left Back    21210 non-null  bool
 34  position_Left Center Back 21210 non-null  bool
 35  position_Left Center Forward 21210 non-null  bool
 36  position_Left Center Midfield 21210 non-null  bool
 37  position_Left Defensive Midfield 21210 non-null  bool
 38  position_Left Midfield 21210 non-null  bool
 39  position_Left Wing    21210 non-null  bool
 40  position_Left Wing Back 21210 non-null  bool
 41  position_Right Attacking Midfield 21210 non-null  bool
 42  position_Right Back   21210 non-null  bool
 43  position_Right Center Back 21210 non-null  bool
 44  position_Right Center Forward 21210 non-null  bool
 45  position_Right Center Midfield 21210 non-null  bool
 46  position_Right Defensive Midfield 21210 non-null  bool
 47  position_Right Midfield 21210 non-null  bool
 48  position_Right Wing   21210 non-null  bool
 49  position_Right Wing Back 21210 non-null  bool
 50  position_Secondary Striker 21210 non-null  bool
 51  period_First Half    21210 non-null  bool
 52  period_Second Half   21210 non-null  bool
 53  period_Third Period  21210 non-null  bool
 54  period_Fourth Period 21210 non-null  bool
```

```

55 period_Penalties           21210 non-null  bool
56 shot_classification_Bottom 21210 non-null  bool
57 shot_classification_Bottom-Corner 21210 non-null  bool
58 shot_classification_Middle   21210 non-null  bool
59 shot_classification_Off-Target 21210 non-null  bool
60 shot_classification_Top     21210 non-null  bool
61 shot_classification_Top-Corner 21210 non-null  bool
dtypes: bool(56), float64(3), int64(3)
memory usage: 2.8 MB

```

```

In [25]: from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE

# Define your features (X) and target (y)
X = cleaned_model_data.drop(columns=['Goal']) # Drop the target column
y = cleaned_model_data['Goal'] # Target column (goal or not)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Check the shapes of the splits
print(X_train.shape, X_test.shape)

smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

```

(16968, 61) (4242, 61)

Logistic Regression

```

In [26]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler

# Feature scaling (Logistic Regression often benefits from scaling)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_resampled) # Use resampled data
X_test_scaled = scaler.transform(X_test)

# Initialize and fit the Logistic Regression model
log_reg = LogisticRegression(random_state=42, max_iter=1000, class_weight='balanced')
log_reg.fit(X_train_scaled, y_train_resampled) # Fit on resampled data

# Predict on the test set
y_pred = log_reg.predict(X_test_scaled)

# Evaluate the model
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

```

Confusion Matrix:

```
[[3572 135]
 [ 295 240]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.96	0.94	3707
1	0.64	0.45	0.53	535
accuracy			0.90	4242
macro avg	0.78	0.71	0.74	4242
weighted avg	0.89	0.90	0.89	4242

RandomForestClassifier

In [28]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Initialize the Random Forest Classifier
model = RandomForestClassifier(random_state=42)

# Train the model on the resampled training data
model.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print(f"Accuracy: {accuracy}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(class_report)
```

Accuracy: 0.9125412541254125

Confusion Matrix:

```
[[3523 184]
 [ 187 348]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.95	0.95	3707
1	0.65	0.65	0.65	535
accuracy			0.91	4242
macro avg	0.80	0.80	0.80	4242
weighted avg	0.91	0.91	0.91	4242

XGBoost

In [30]:

```
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from xgboost import XGBClassifier
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_curve
from imblearn.over_sampling import SMOTE

# Apply SMOTE to oversample the minority class
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

# Define the model with class weights (scale_pos_weight)
xxgb_clf = XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    max_depth=6,
    learning_rate=0.1,
    n_estimators=100,
    scale_pos_weight=len(y_train_resampled) / sum(y_train_resampled), # Adding
    use_label_encoder=False
)

# Train the model
xxgb_clf.fit(X_train_resampled, y_train_resampled)

# Get probability scores
yy_prob_xgb = xxgb_clf.predict_proba(X_test)[:, 1]

# Compute Precision-Recall Curve
precision, recall, thresholds = precision_recall_curve(y_test, yy_prob_xgb)

# Plot the Precision-Recall curve
plt.plot(thresholds, precision[:-1], label="Precision", color="blue")
plt.plot(thresholds, recall[:-1], label="Recall", color="orange")
plt.xlabel("Threshold")
plt.ylabel("Score")
plt.legend()
plt.title("Precision-Recall Tradeoff (Goal Prediction)")
plt.grid()
plt.show()

# Set the threshold lower to maximize recall for class 1 (goals)
# For example, you can choose a threshold value where recall is higher but precision is lower
selected_threshold = 0.65 # Adjust the threshold based on the curve


# Apply the new threshold
yy_pred_xgb_adjusted = (yy_prob_xgb >= selected_threshold).astype(int)

# Evaluate the model with the adjusted threshold
conf_matrix = confusion_matrix(y_test, yy_pred_xgb_adjusted)
print("Confusion Matrix (Adjusted Threshold - Prioritize Goals):\n", conf_matrix)

# Plot confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix, display_labels=[0, 1])
```


```
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix (Adjusted Threshold - Goals)')
plt.show()

# Print classification report
print("Classification Report (Adjusted Threshold - Prioritize Recall for Goals):")
```



Confusion Matrix (Adjusted Threshold - Prioritize Goals):

```
[[3517  190]
 [ 99  436]]
```



Classification Report (Adjusted Threshold - Prioritize Recall for Goals):

	precision	recall	f1-score	support
0	0.97	0.95	0.96	3707
1	0.70	0.81	0.75	535
accuracy			0.93	4242
macro avg	0.83	0.88	0.86	4242
weighted avg	0.94	0.93	0.93	4242

Catboost

```
In [31]: from catboost import CatBoostClassifier

# Apply SMOTE to the training data to handle class imbalance
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

# Define and train the CatBoost model
model = CatBoostClassifier(
    iterations=500,
    learning_rate=0.1,
    depth=10,
    loss_function='Logloss',
    eval_metric='Accuracy',
    verbose=100
)

model.fit(X_train, y_train, eval_set=(X_test, y_test), early_stopping_rounds=10)
```

```
# Predict and evaluate
y_pred = model.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0:      learn: 0.8952735      test: 0.8903819 best: 0.8903819 (0)      total: 17
5ms      remaining: 1m 27s
Stopped by overfitting detector (10 iterations wait)
```

```
bestTest = 0.9398868458
```

```
bestIteration = 55
```

```
Shrink model to first 56 iterations.
```

```
[[3609  98]
```

```
[ 157 378]]
```

	precision	recall	f1-score	support
0	0.96	0.97	0.97	3707
1	0.79	0.71	0.75	535
accuracy			0.94	4242
macro avg	0.88	0.84	0.86	4242
weighted avg	0.94	0.94	0.94	4242

In []: